

Memory Virtualization:

Time Sharing, Base+Bounds, Segmentation

OSTEP Chapters 13, (14), 15, 16:

<http://pages.cs.wisc.edu/~remzi/OSTEP/vm-intro.pdf>

<http://pages.cs.wisc.edu/~remzi/OSTEP/vm-api.pdf>

<http://pages.cs.wisc.edu/~remzi/OSTEP/vm-mechanism.pdf>

<http://pages.cs.wisc.edu/~remzi/OSTEP/vm-segmentation.pdf>

Jan Reineke

Universität des Saarlandes

Virtualization

Virtual CPU:


illusion of private CPU registers

Virtual RAM:

illusion of private memory

Motivation for Virtualization

Uniprogramming: One process runs at a time



Disadvantages:

- Only one process runs at a time
- Process can destroy OS

Memory Virtualization Goals

Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

Protection

- Integrity: Cannot corrupt OS or other processes
- Privacy/confidentiality: Cannot read data of other processes

Efficiency

- Do not waste memory resources (minimize fragmentation)

Sharing

- Cooperating processes can share portions of address space


Abstraction: Address space

Address space:

Each process has set of addresses that map to bytes

Address space has *static* and *dynamic* components:

- Static: Code and global variables
- Dynamic: Stack and Heap



Motivation for dynamic memory

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time:
 - often depends on program inputs
- Would have to statically allocate memory for the “worst case” → inefficient

Examples of dynamic memory allocation

Examples:

- Recursive procedures
- Complex data structures:
lists, trees, hash maps, etc.


Two types of dynamic allocation:

- Stack
- Heap

Stack organization

Definition: Memory is freed in opposite order from allocation:

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```



Simple and efficient implementation:
Pointer separates allocated and freed space

- *Allocation:* Decrement pointer
- *Deallocation:* Increment pointer

Where are stacks used?

OS uses stack for procedure call frames:
local variables + parameters on the stack

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```

Heap organization

Definition: Allocate from any random location: malloc(), new()

- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

Advantage:

- Works for all data structures

Disadvantages:

- Allocation more complex, can be slow
- Fragmentation

Divison of work: OS + library

- OS gives big chunk of free memory to process;
library manages individual allocations



Quiz: Match that address allocation

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

Address	Location
x	Static data
main	Code
y	Stack
z	Stack
*z	Heap

Memory accesses

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```



```
movl    -0x14(%rbp), %edi
addl    $0x3, %edi
movl    %edi, -0x14(%rbp)
```

(x86 Assembler)
%rbp is pointing to the base
of the current stack frame

```
gcc -o exp exp.c
otool -tv exp
(or objdump under Linux)
```

Quiz: Memory accesses?

Initial:

%rip = 0x10 (PC)

%rbp = 0x200 (Basis des Stacks)

1a) Instruction fetch at address 0x10

1b) Load from address $0x200 + -0x14 = 0x1EC$

2a) Instruction fetch from address 0x13

2b) No memory access

3a) Instruction fetch at address 0x19

3b) Store to address 0x1EC

 **0x10: movl -0x14(%rbp), %edi**

 **0x13: addl \$0x3, %edi**

 **0x19: movl %edi, -0x14(%rbp)**

How to virtualize memory?

Problem: How to run multiple processes simultaneously?

Challenge: Addresses are “hardcoded” into process binaries

Possible solutions for mechanisms:

1. Time sharing
2. Static relocation
3. Dynamic relocation
4. Segmentation

1. Time sharing of memory

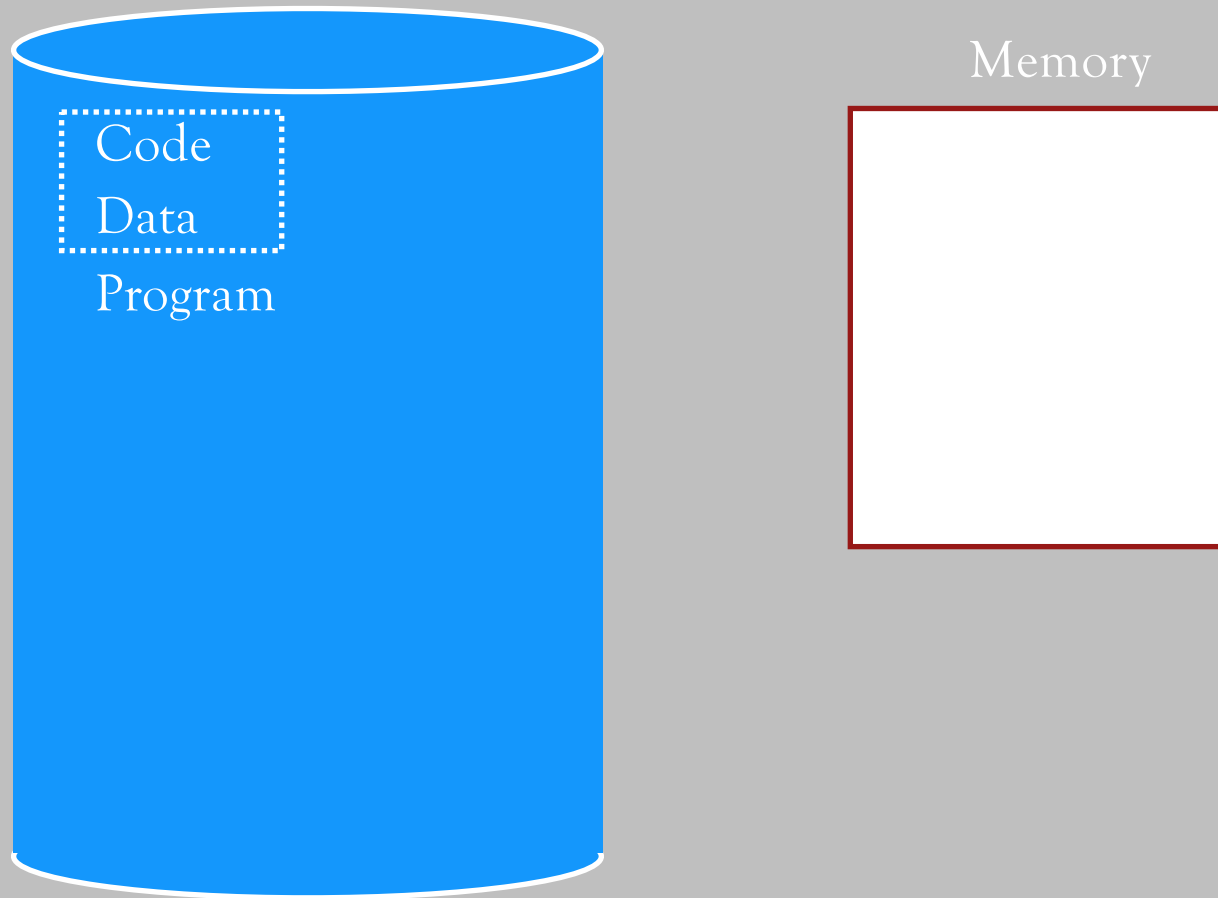
Try similar approach to how OS virtualizes CPU

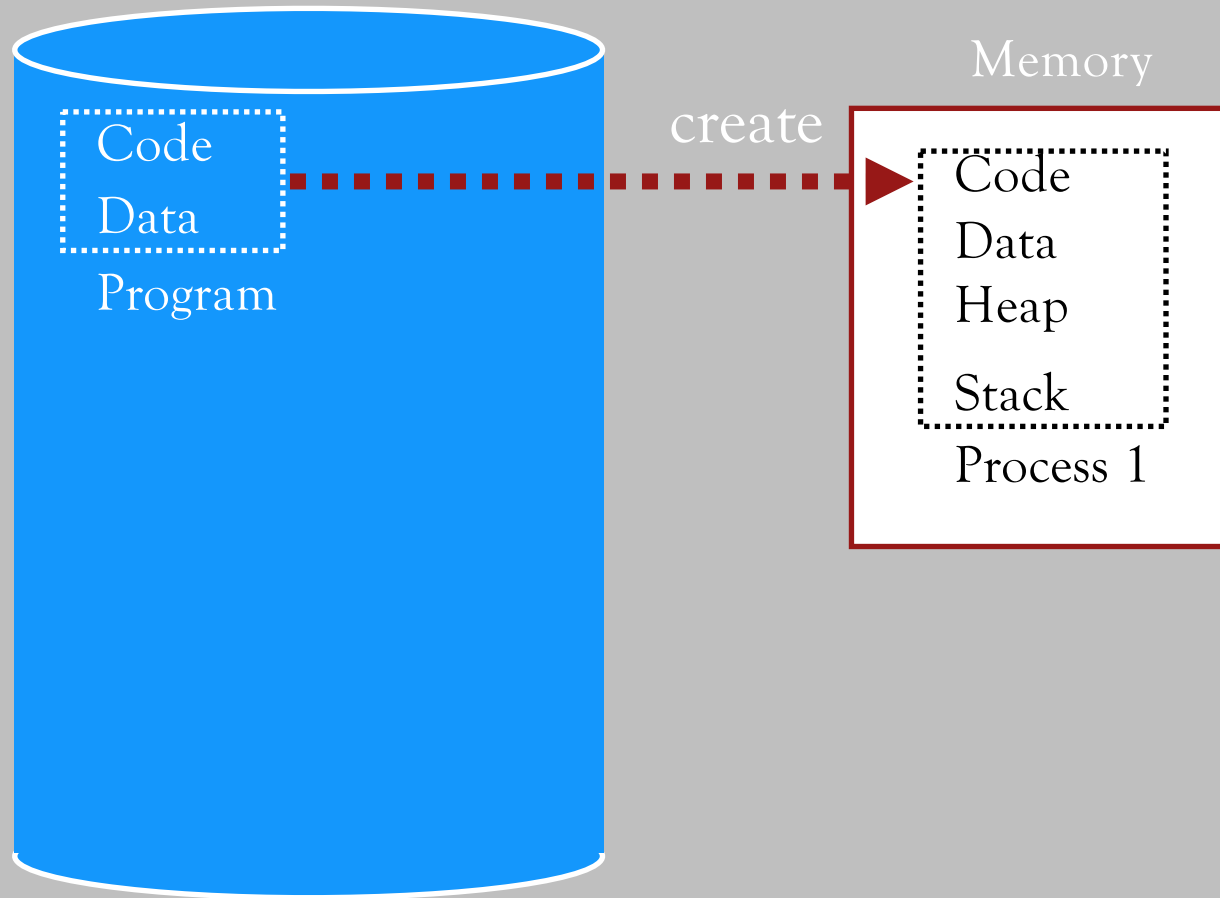
Observation: OS gives illusion of private CPUs by saving CPU registers to memory when a process isn't running



Approach:

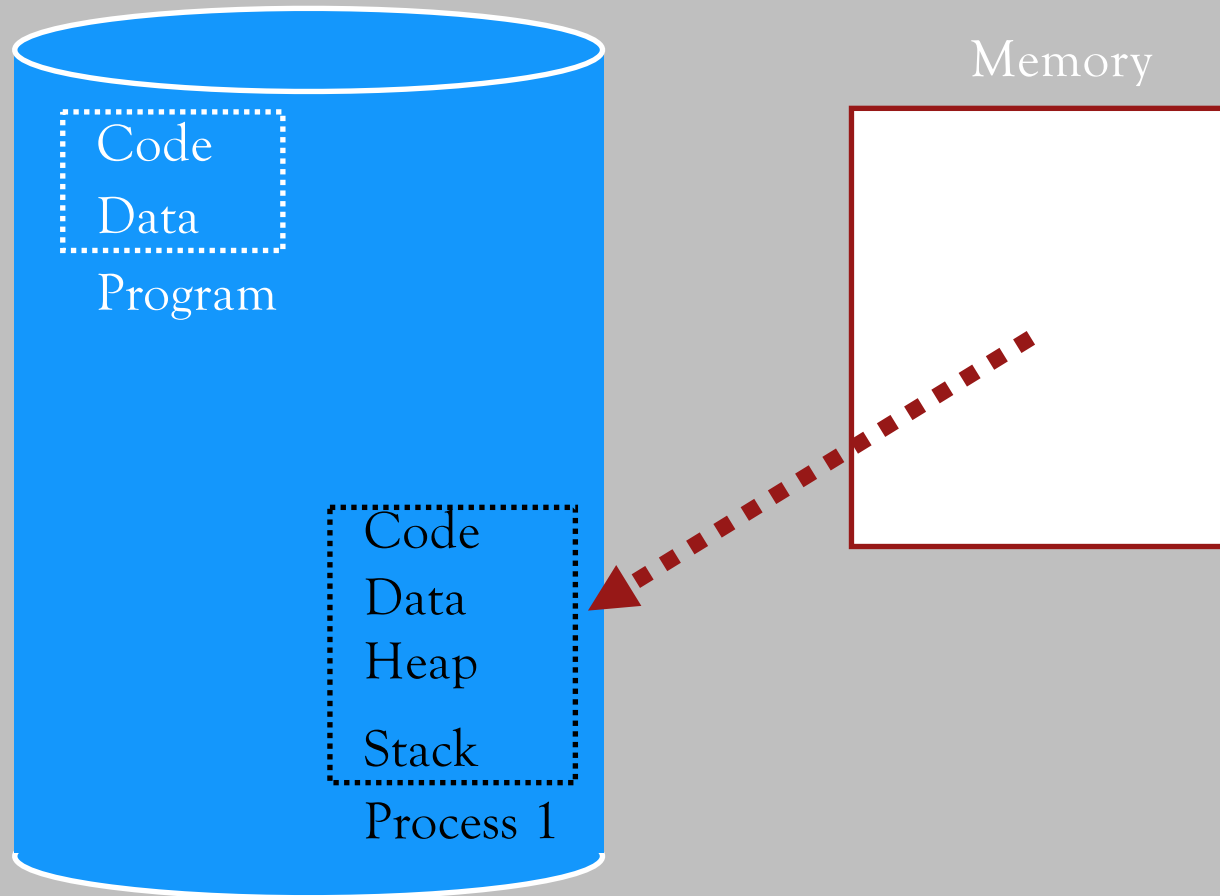
Save memory contents to disk when process isn't running

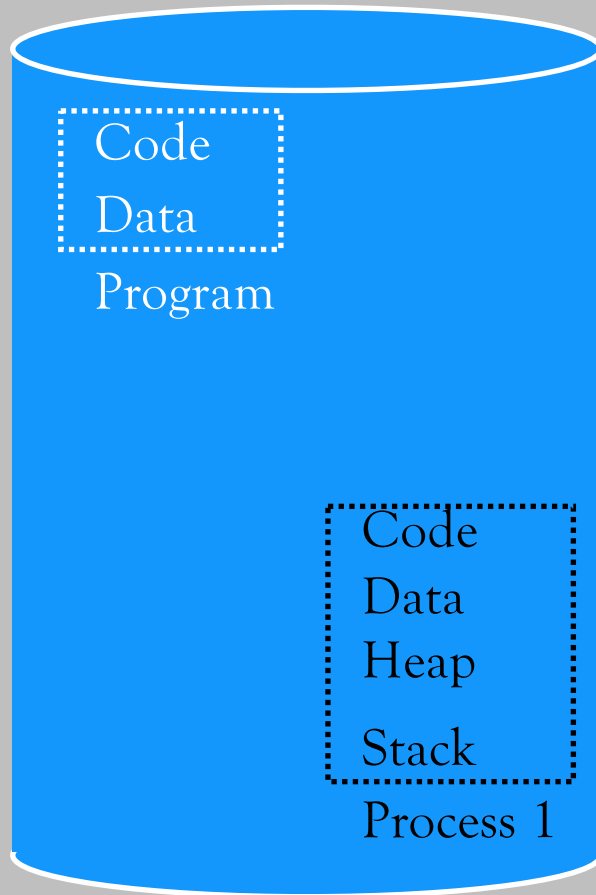
Example: Time sharing





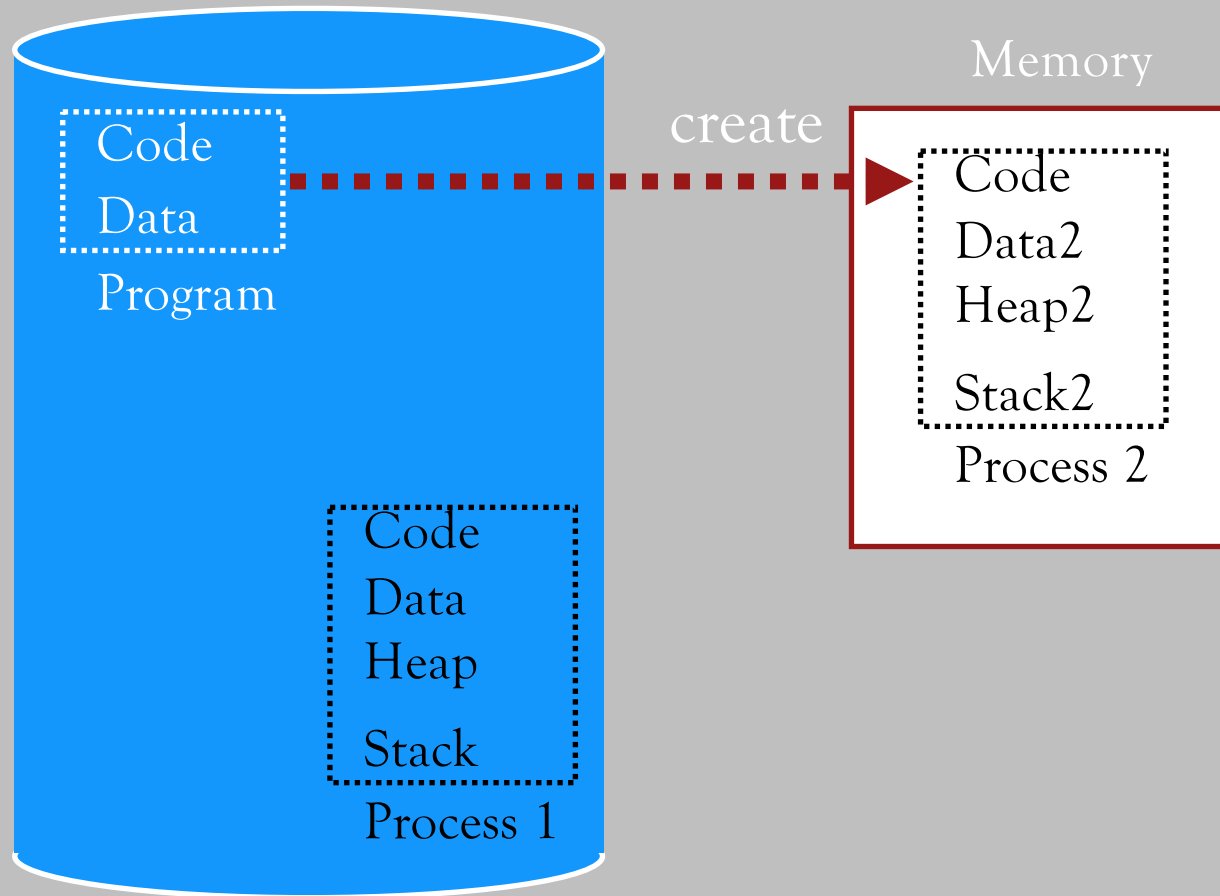


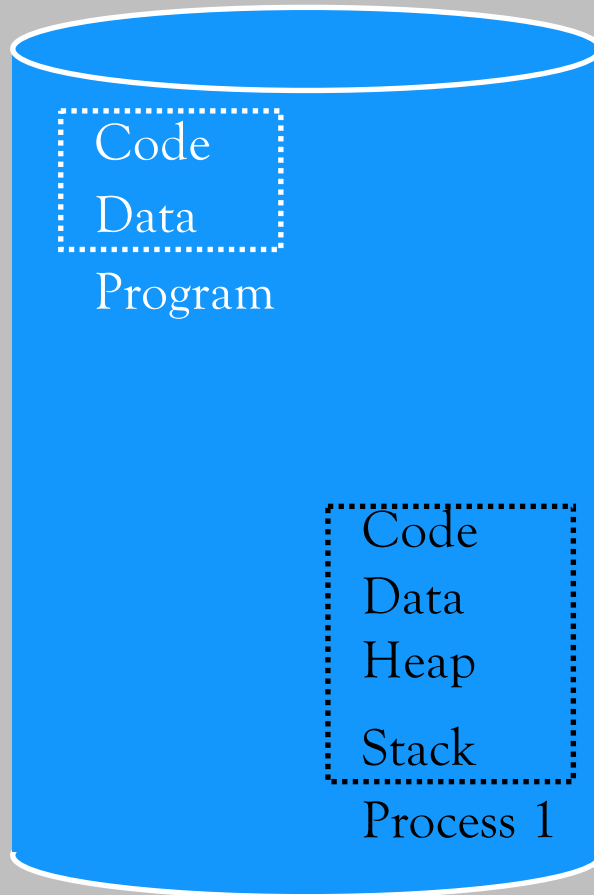





Memory

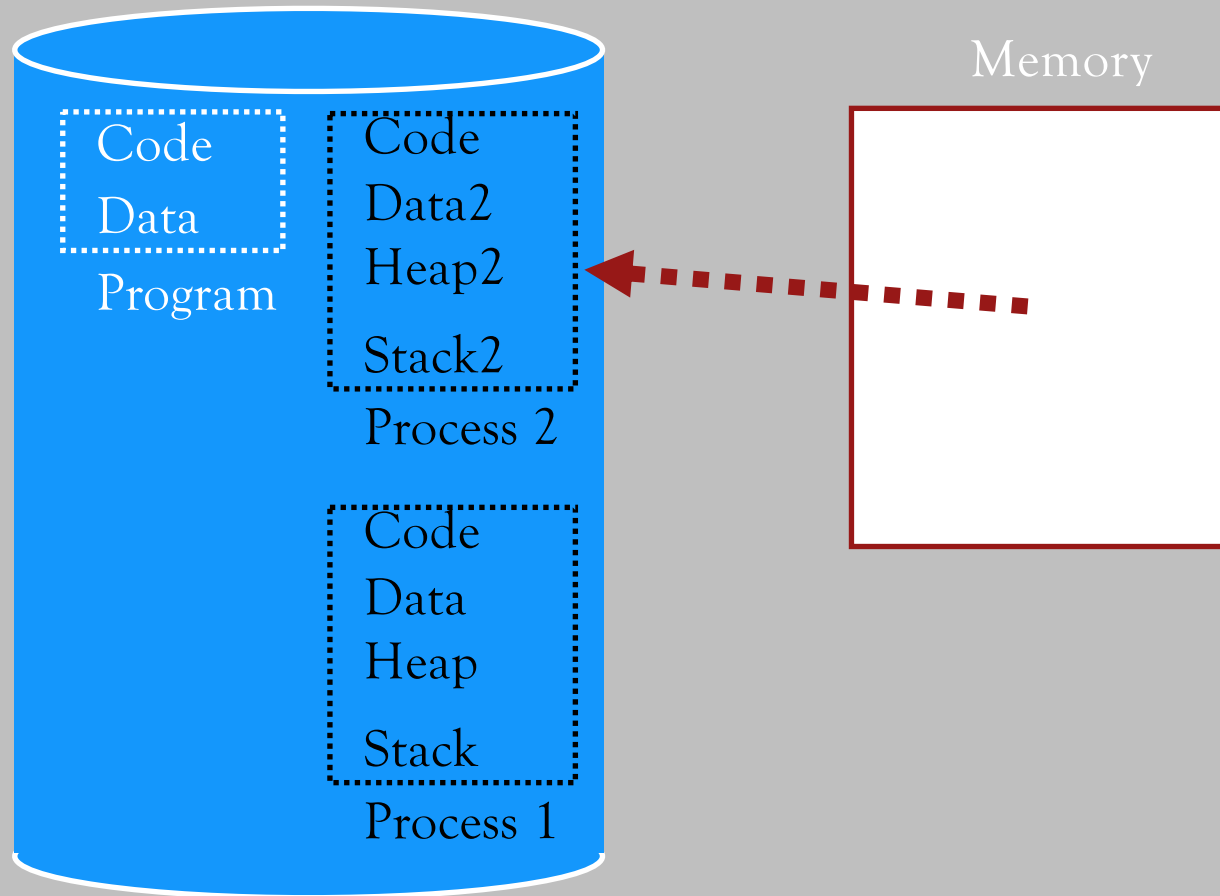


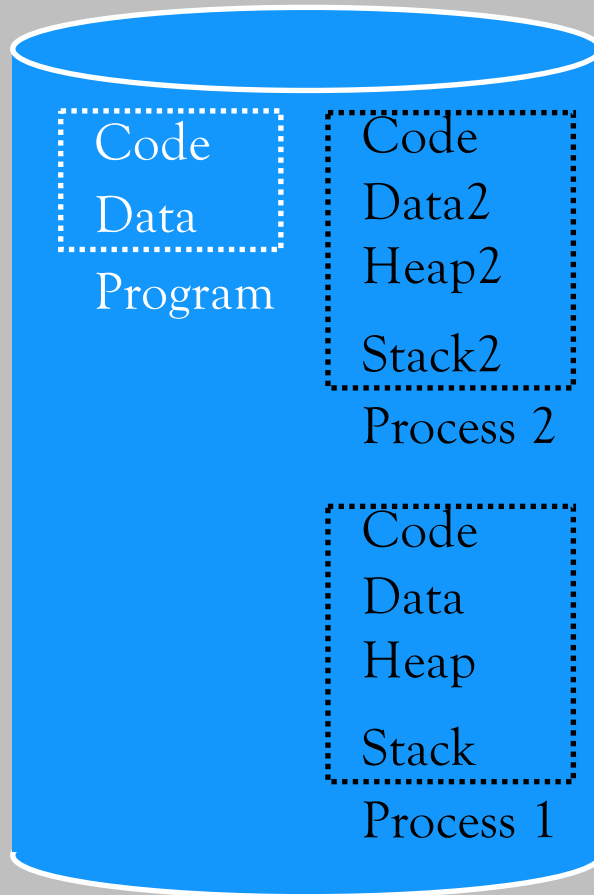




Memory

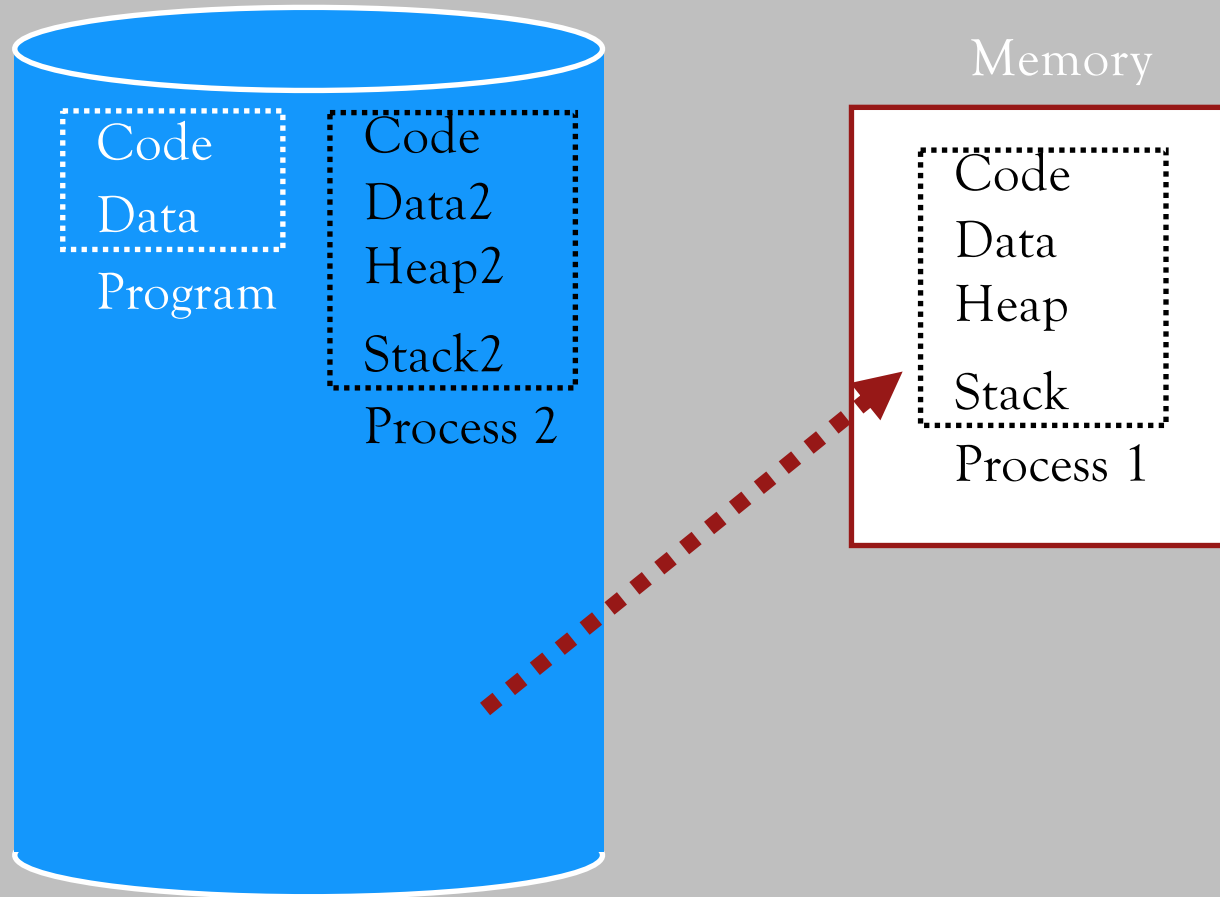








Memory







Memory



Problem?

Problems with time sharing memory

Problem:

Extremely poor performance, as copying expensive

Better alternative: **space sharing**

- Physical memory is divided across several processes

2. Static relocation

Idea: OS **rewrites** each program before loading it as a process in memory:

- Each rewrite for different process uses **different** addresses and pointers
- Change jumps, loads of static data

Example: Static relocation

- 0x10: movl 0x8(%rbp), %edi
- 0x13: addl \$0x3, %edi
- 0x19: movl %edi, 0x8(%rbp)

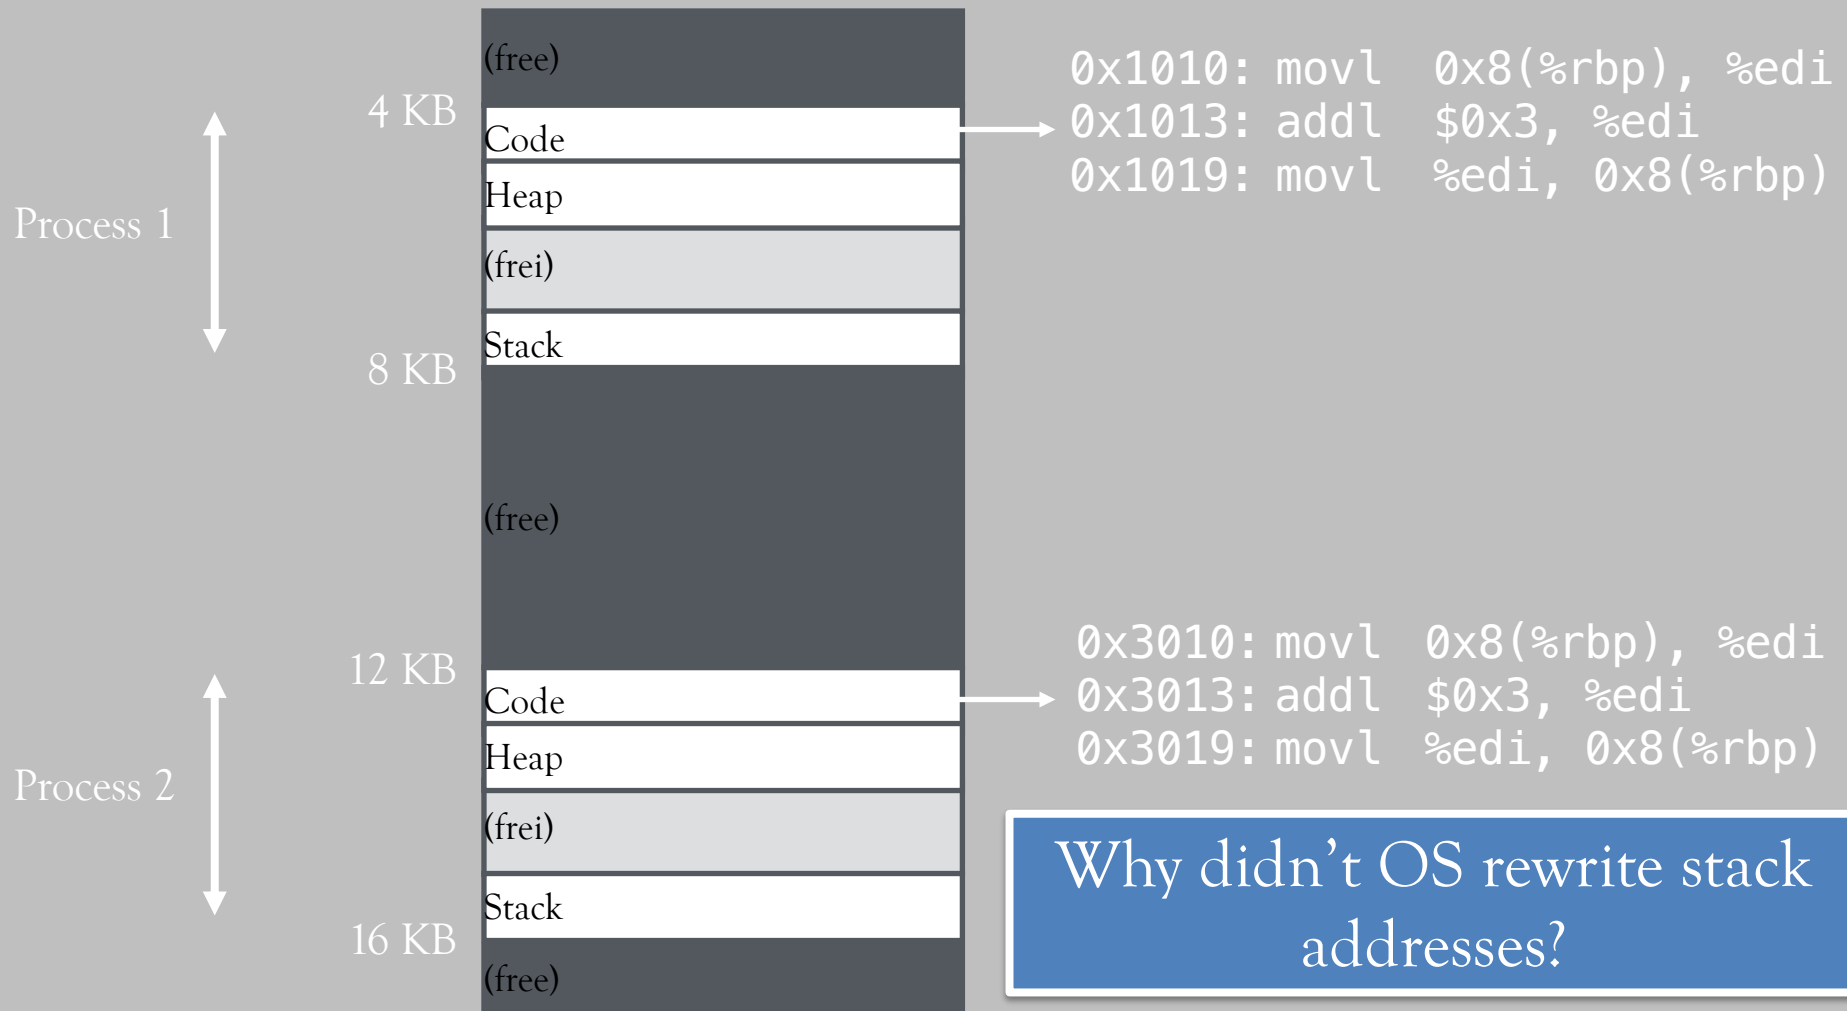
rewrite

0x1010: movl 0x8(%rbp), %edi
0x1013: addl \$0x3, %edi
0x1019: movl %edi, 0x8(%rbp)

rewrite

0x3010: movl 0x8(%rbp), %edi
0x3013: addl \$0x3, %edi
0x3019: movl %edi, 0x8(%rbp)

Example: Static relocation



Static relocation: Disadvantages

No protection:

Process can destroy (and spy on) OS
or other processes

Cannot move address space after it has been placed
→ possible fragmentation


3. Dynamic relocation

Goals:

1. Allow **relocation** of processes even **after** they have been started
2. **Protection** of processes from one another

3. Dynamic relocation

- Requires hardware support:
Memory Management Unit (MMU)
- MMU dynamically changes process address at every memory access
 - Process generates **virtual (or logical) address** (in their address space)
 - Hardware uses **physical (or real) addresses**



Hardware support for dynamic relocation

Two operating modes:

- **Kernel (protected, privileged) mode:** reserved for OS
 - Can manipulate contents of MMU
 - Allows OS to access all of physical memory
- **User mode:** for user processes
 - MMU translates virtual addresses to physical addresses

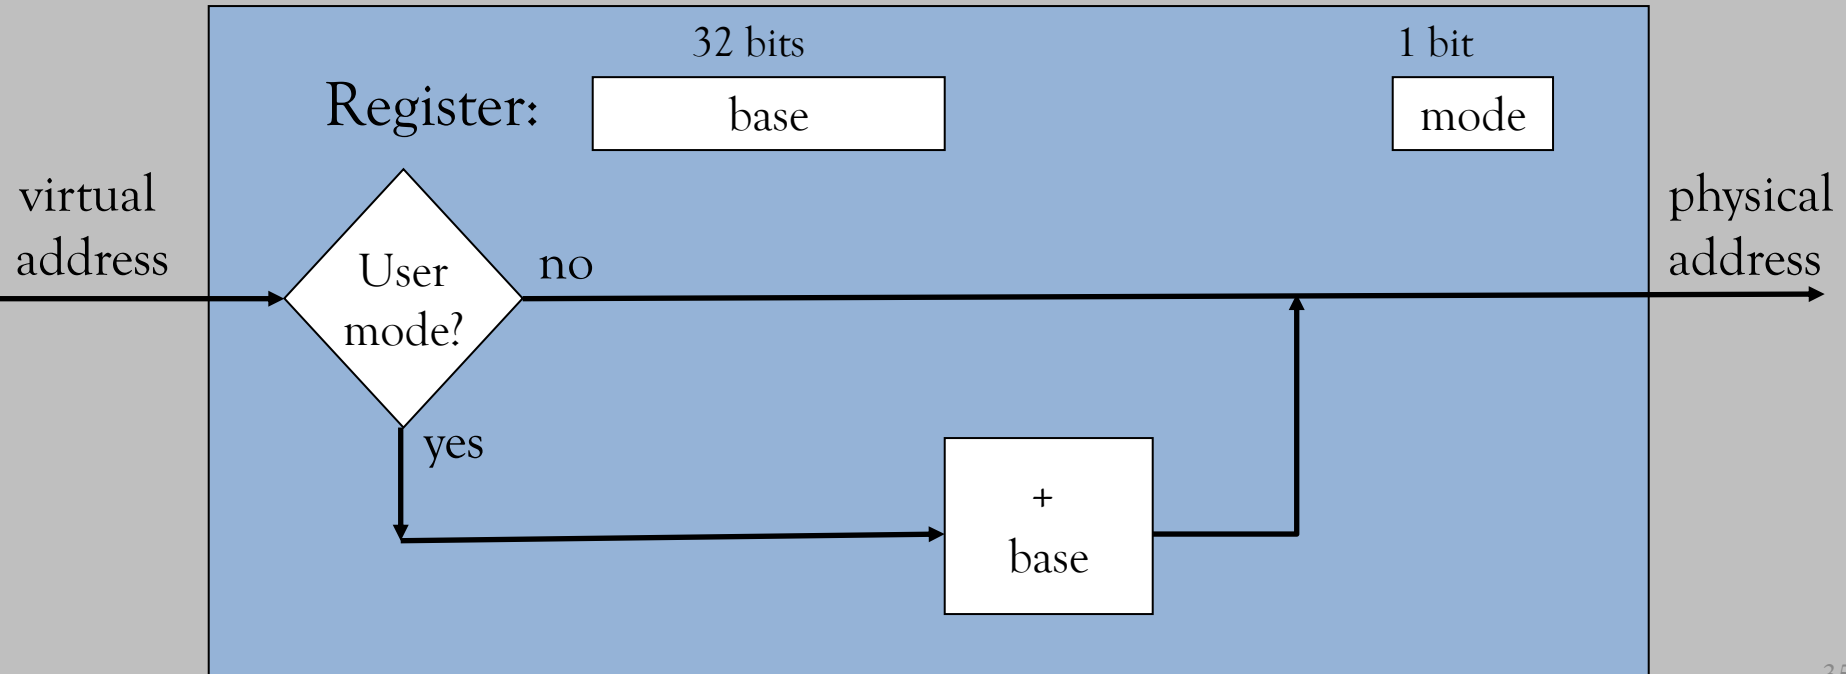
Minimal MMU contains **base register** for translation

- **base:** start location for address space

Implementation of dynamic relocation: Base register


MMU sums value of base register onto virtual addresses to obtain physical addresses

Memory Management Unit




Dynamic relocation with base register

OS writes **correct value** to **base register** upon each context switch



Example: **Dynamic relocation**



P1 running



P2 is running



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1

Quiz: Who controls the base register?

What entity **does translation** of addresses?

(a) user process, (b) OS, or (c) HW

Which entity **modifies** the base register?

(a) user process, (b) OS, or (c) HW

Quiz: Who controls the base register?

What entity **does translation** of addresses?

(a) user process, (b) OS, or (c) HW

Which entity **modifies** the base register?

(a) user process, (b) OS, or (c) HW

		Virtuell	Physisch
0 KB			
1 KB	P1	P1: load 100, R1	load 1124, R1
2 KB		P2: load 100, R1	load 4196, R1
3 KB		P2: load 1000, R1	load 5096, R1
4 KB	P2	P1: load 1000, R1	load 2024, R1
5 KB		P1: store 3072, R1	store 4096, R1
6 KB			

Can **P1** modify data of **P2**?

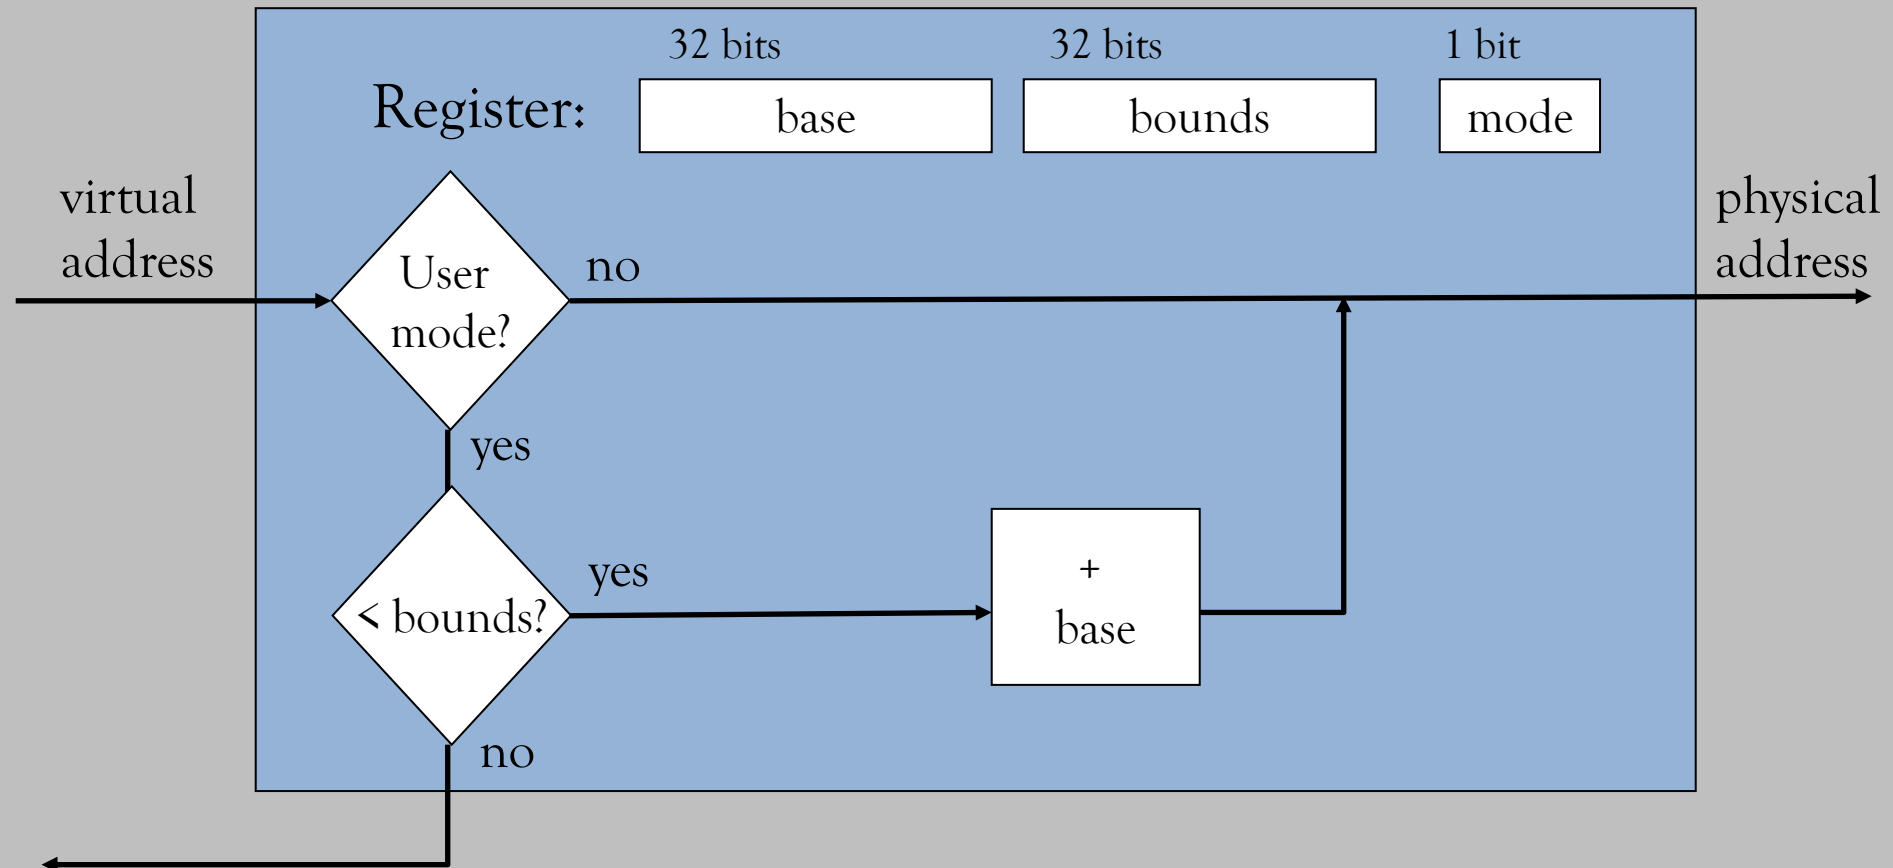
Dynamic relocation with base + bounds


Idea: Limit address space with a “bounds register”

- **Base register:** smallest physical address
- **Bounds register:**
Size of the process's virtual address space
- OS **kills** process if process loads/stores beyond bounds

Implementation of base + bounds

Memory Management Unit



		Virtual	Physical
0 KB			
1 KB	P1 	P1: load 100, R1	load 1124, R1
2 KB		P2: load 100, R1	load 4196, R1
3 KB		P2: load 1000, R1	load 5096, R1
4 KB	P2	P1: load 100, R1	load 2024, R1
5 KB		P1: store 3072, R1	interrupt OS!
6 KB			

Extension of OS for dynamic relocation

Add fields for base and bounds of address space in process control blocks (PCB)

Context switch (in kernel mode):

1. Load base and bounds values of new process from PCB into MMU registers
2. Switch to user mode and jump to new process

Precondition for security:

User processes are unable to modify base and bounds registers


→ Ensured by execution in user mode

Advantages of dynamic relocation

1. Supports dynamic relocation of processes at **runtime**
 2. Provides **protection** across address spaces
- Simple and cheap to implement:
few registers, little logic in MMU
 - *Fast*: Add and compare can be performed in parallel

Disadvantages of dynamic relocation

- Each process must be allocated **contiguously** in physical memory
 - *Internal* fragmentation
 - *External* fragmentation
- **No partial sharing:**
Cannot share
limited parts of address space



4. Segmentation

Divide address space into **logical segments**:

Code, Stack, Heap

Each segment can **independently**:

- be placed in physical memory
- grow and shrink
- be protected
(separate read/write/execute protection bits)

Segmented addressing

How does process designate a particular segment?

- Use part of virtual address:
 - most-significant bits select segment
 - other bits encode offset within segment

Segmentation: Implementation

MMU contains **segment table** (per process):

- Each segment has own base and bounds, protection bits
- *Example:* 14-bit virtual address, 4 segments
 - How many segment bits? 2
 - How many for offset? 12

Segment	Base	Bounds	R	W
0	0x2000	0x6fff	1	0
1	0x0000	0x4fff	1	1
2	0x3000	0xffff	1	1
3	0x0000	0x0000	0	0

Quiz: Address translations with segmentation

Segment	Base	Bounds	R W
0	0x2 <u>000</u>	0x <u>6ff</u>	1 0
1	0x0 <u>000</u>	0x4ff	1 1
2	0x3 <u>000</u>	0xffff	1 1
3	0x0 <u>000</u>	0x000	0 0

Translate logical addresses (in hex) to physical addresses?

0x0240: 0x2240

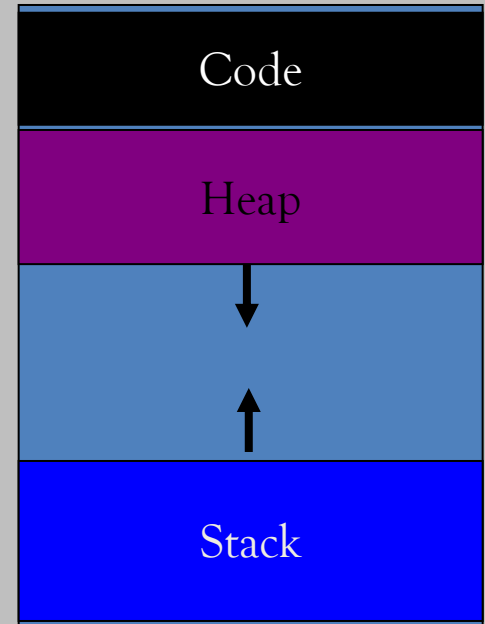
0x1108: 0x0708

0x265c: 0x3b5c

0x3002:

Advantages of segmentation

- Enables “sparse” allocation of address space:
 - Stack and heap grow independently of each other
→ no **internal** fragmentation
 - *Heap*: If no data on free list, dynamic memory allocator (in library) requests more from OS (e.g., UNIX malloc calls sbrk())
 - *Stack*: OS recognizes references outside legal segment, extends stack implicitly
- Different protection for different segments:
 - E.g. read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



Disadvantages of segmentation

Each segment must be allocated **contiguously** in memory

- **External** fragmentation of the physical memory
- Paging as solution next lecture

Conclusion

- HW + OS work together to virtualize memory
 - Memory Management Unit supports fast address translation in HW
 - OS only involved upon context switches or errors